

HMM

Since the MILK command language tends to be much more verbose than the recipes themselves, it is essential for us to allow multiple, related MILK commands to map into a single sentence. In order to do so, we developed a HMM-like model to segment the sequential commands into coherent sentence groups. In our model, the hidden variables, y_1, \dots, y_n can take on values $1, 2, \dots, k$ representing the number of MILK commands that generate a single sentence. The visible variables x_1, \dots, x_n are sequences of commands up to length k . To simplify matters and avoid sparsity issues, we simply concern ourselves with the MILK commands, disregarding their parameters. Furthermore we ignore all *create_ing* and *create_tool* commands, since each of those commands forms a sentence. In the vast majority of cases, no more than 4 MILK commands correspond to a single sentence. To generate English, we first tackled the problem of segmenting the sequential MILK command into coherent groups, so that multiple, related MILK commands may generate a single sentence. An illustration of our model can be seen in Figure 1 below.

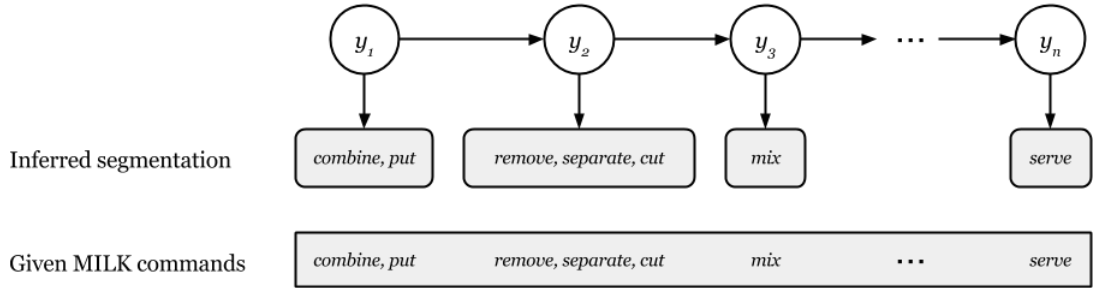


Figure 1: Graphical model for the HMM-variant

In our hidden Markov model, the MILK commands segmentation reduces to a sequence labeling problem. The input is a sequence of MILK commands of length n : x_1, \dots, x_n . The output is another sequence y_1, \dots, y_n representing the generated sentence groups.

During training time, we estimate the HMM parameters σ and τ by maximum likelihood estimation based on our preprocessed training recipes, each of which is a sequence of MILK commands with their group labels:

$$\hat{\sigma}_{y,y'} = \frac{n_{y,y'}(\mathbf{y})}{n_{y,o}(\mathbf{y})}, \quad \hat{\tau}_{y,x} = \frac{n_{y,x}(\mathbf{x},\mathbf{y})}{n_{y,o}(\mathbf{x},\mathbf{y})}$$

where, $n_{y,y'}$ is the number of times label y' followed by y and $n_{y,x}$ is the number of times x is labeled as y . To deal with the sparse data problem, we compute the smoothed $\tilde{\tau}_{y,x}$ by adding a pseudo-count α :

$$\tilde{\tau}_{y,x} = \frac{n_{y,x}(\mathbf{x},\mathbf{y}) + \alpha}{n_{y,o}(\mathbf{x},\mathbf{y}) + \alpha|W|}$$

where $|W|$ represents the total count of MILK commands in the training data. We perform a line search based on our held-out data to choose the best α among a set of candidate values.

During testing time, given the estimated $\hat{\sigma}$, $\tilde{\tau}$ parameters and a sequence of MILK commands, we perform the Viterbi algorithm to find the most likely label sequence \mathbf{y} . In fact, we adapt the algorithm to find the top- k most likely sequences of labels. Finally, we recover the MILK command segmentation based on the labels.

To evaluate our model, we compared the Viterbi segmentations to the true ones by treating every boundary between MILK commands as an individual binary classification problem. Finally, we considered the F-score of our model combined over all sentences in our test set. Our average F-score was on average 0.766 for $\alpha = 10^{-5}$.

This is the first part in our MILK-to-English pipeline. The segmentation result is to be used in a “just-in-time” way: As we go through a sequence of MILK commands, at command x_i , we wish to know the distribution over the number of commands that follow x_i and should be grouped together with x_i to form a sentence. We compute these distributions by simply counting the corresponding frequencies in our k best sequences.